



1

## Acknowledgement

- These slides are based on the following page:
- <https://savan77.github.io/2017-04-19-ml-part1/>
- **Why?**
  - They provide the code that you can try out.

Yaay!!

2

## Supervised Learning

- Supervised learning problems can be divided into two types:
  - Regression, and
  - Classification

3

## Linear Regression

- Example:** Property prices have become a hot topic and we happen to have a dataset about houses:

area (square feet)	price (1k\$)
3456	600
2089	395
1416	232
...	...

(the data happen to be for houses in San Francisco – measuring the house area in square feet and house price in USD) – the example code uses the a similar dataset for the Boston area in the US

- Say, we try to predict house prices based on their areas. Then, **price** becomes our **target variable** (**y**), and **area** is our only **input variable** (**x**)

4

## Linear Regression

- The **Regression Problem** is the problem of learning the **target function**  $f(x)$  that outputs the target value  $y$  for a given input  $x$ . In our example,  $\text{price} = f(\text{area})$
- If we believe that linear regression can be a good model for the target function  $f$ ,  $f$  would take the form of a linear function:  $y = f(x) = a * x + b$


5

## Linear Regression

- For our dataset:

area (square feet)	price (1k\$s)
3456	600
2089	395
1416	232
...	...


- If we visualise it:



6

## Linear Regression

- If we choose a linear function to model the **target function**  $\text{price} = f(\text{area}) = \mathbf{a} * \text{area} + \mathbf{b}$ ; the question is: **What are the “good” values for  $\mathbf{a}$  and  $\mathbf{b}$ ??**
- Different values of  $\mathbf{a}$  and  $\mathbf{b}$  will give us different lines:



- The aim of ML algorithms is to **search** for the **good** values for the **parameters  $\mathbf{a}$  and  $\mathbf{b}$**  (**optimization** = minimizing the **error**)

7

## Linear Regression

- Instead of writing  $\mathbf{y} = \mathbf{f}(\mathbf{x}) = \mathbf{a} * \mathbf{x} + \mathbf{b}$ , we can adopt the following standardized form:  $\mathbf{y} = \mathbf{f}(\mathbf{x}) = \mathbf{w}_0 + \mathbf{w}_1 * \mathbf{x}$  (replacing the parameters  $\mathbf{a}$  by  $\mathbf{w}_1$  and  $\mathbf{b}$  by  $\mathbf{w}_0$ )

(this is why parameters are also called “**weights**”)

- More generally, if instead of having just one input variable  $\mathbf{x}$ , we have  $m$  input variables  $\mathbf{x}_1, \dots, \mathbf{x}_m$  then the linear regression target function will be written as:

$$\mathbf{y} = \mathbf{f}(\mathbf{x}_1, \dots, \mathbf{x}_m) = \mathbf{w}_0 + \mathbf{w}_1 * \mathbf{x}_1 + \dots + \mathbf{w}_m * \mathbf{x}_m$$

8

## Linear Regression

- Then, given a model represented by the parameters  $W = (w_0, w_1, \dots, w_m)$ , we can calculate the error of this model using the **cost function**  $J(W)$ :

$$J = \frac{1}{n} \sum_{i=1}^n (pred_i - y_i)^2$$



where:

- $n$  is the number of data points in the training dataset.
- For the  $i^{\text{th}}$  data point  $(x_1^i, \dots, x_m^i, y_i)$ :  $pred_i = f(x_1^i, \dots, x_m^i)$  is the model prediction.
- The above cost function  $J$  is the mean squared error.

9

## Gradient descent

- Now that we have a way to measure the model's error in the cost function  $J(W)$ , we want to find the value of weights for which error is smallest (**minimizing** cost function).
- Gradient Descent** is one of the most popular and widely used optimization algorithm.



10

## More terminologies


- Cost function is also called “loss function”  
Model loss  $\sim$  model error
- There are many variants of Gradient Descent (GD):
  - **Batch GD**: cost function is calculated on the entire training data ( $n$  = size of dataset)  $\rightarrow$  slow
  - **Stochastic Gradient Descent (SGD)**: cost function is calculated on each training example ( $n = 1$ )
  - **Mini-batch gradient descent**: cost function is calculated on every mini-batch of  $k$  training examples ( $k <$  size of dataset)
- For more info: <https://ruder.io/optimizing-gradient-descent/index.html#gradientdescentvariants>

11

## More terminologies

- What if a straight line (or hyperplane in the case of multiple input variables) does not fit the data well?
- Linear regression will still do the job by introducing a new feature  $x_{\text{new}} = x^2$
- The target function will then look like this:

$$y = f(x) = w_0 + w_1 * x + w_2 * x^2$$




- And of course, we can have other features for  $x^3$ ,  $x^4$ , etc.

12

## Overfitting

- If we are overzealous in fitting the training data (e.g. by introducing  $x^3$ ,  $x^4$ , etc.) we may end up **overfitting** the data:



- The above model is unlikely to perform well on unseen data, especially around the **anomaly data points**.
- But, *the aim of ML models is to perform well on unseen data.*

13

## Underfitting

- The opposite of **overfitting** is **underfitting**:
  - Typically that is when the model is unable to capture the relationship between the input variables and the target variables accurately, generating a high error rate on both the training set and unseen data.
  - Perhaps we chose the wrong representation?
    - E.g. we should use decision tree (or SVM, or kNN, or Naïve Bayes) instead of linear regression??
  - Perhaps there are no patterns that can generalise for the data at hand??

14

## Overfitting - How to deal with **noise**?

- Instead of

$$J(W) = \frac{1}{n} \sum_{i=1}^n (w_0 + w_1 x_1^i + \dots + w_m x_m^i - y_i)^2$$

- We will add the regularization term:

$$J(W) = \frac{1}{n} \left( \sum_{i=1}^n (w_0 + w_1 x_1^i + \dots + w_m x_m^i - y_i)^2 + \lambda \sum_{j=1}^m w_j^2 \right)$$


**Regularization term  $L_2$**

When  $\lambda$  is sufficiently large,  
it will diminish the impact of  
the training examples

15

## What about logistic regression?

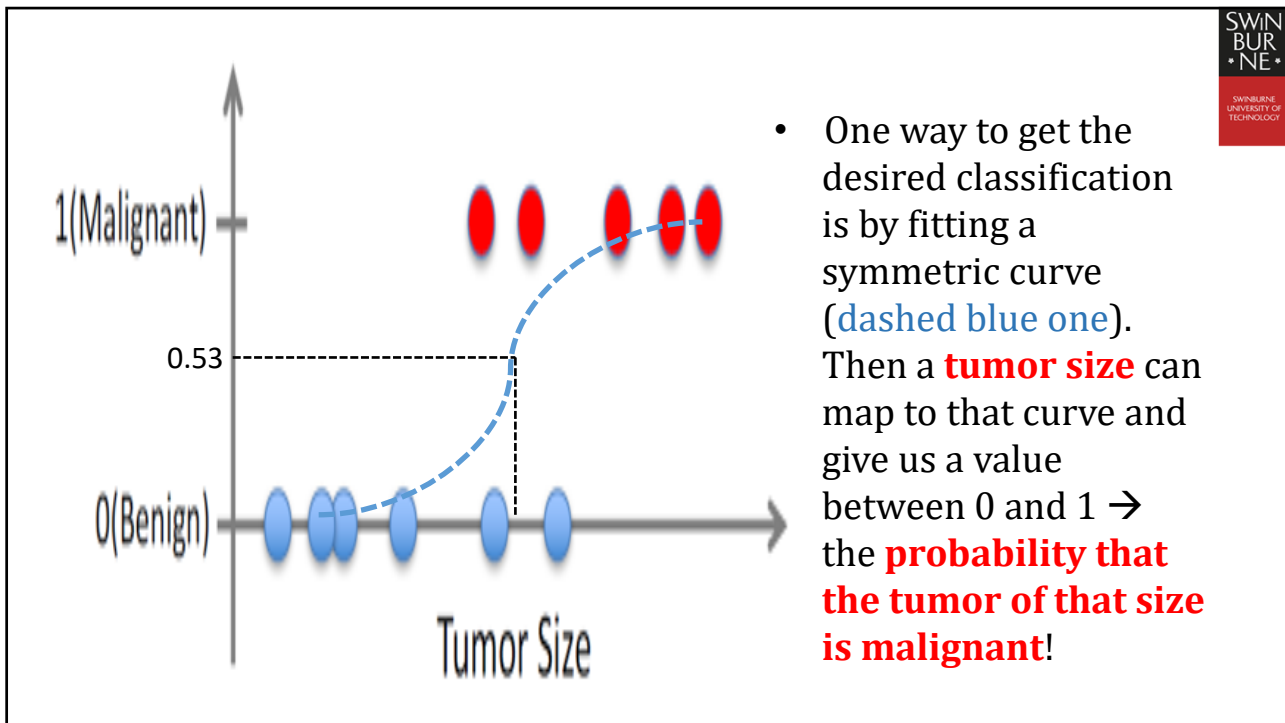
- It is actually a classification technique! How??
- Remember this classification problem?



Based on example by Andrew Ng

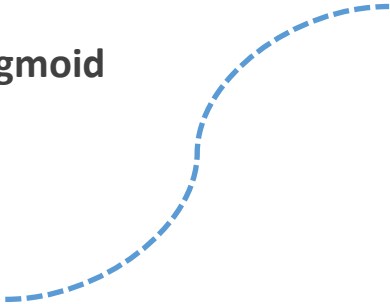
16






17

- A function with this shape is called **logistic function** or **sigmoid function**:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$


- As we don't want to return the probability (requiring some interpretation), what we can do is to map this probability outcome to the classification:
  - when the **prob  $\geq 0.5$** , the tumor will be *classified by our model* as **Malignant**; **else** it will be classified as **Benign**.

18




19

## Logistic regression - mathematically

- $y \sim f(x) = \frac{1}{1+e^{-g(x)}} = \sigma(g(x))$
  - (Well, more precisely  $\text{Prob}(y=1 | x) = \frac{1}{1+e^{-g(x)}}$ )
  - E.g.,  $\text{Prob}(y = \text{Malignant} | \text{size}) = \frac{1}{1+e^{-g(\text{size})}}$
  - and  $g(x) = w_0 + w_1 * x$
  - Of course, for multivariate problem:
- $$g_w(x_1, \dots, x_m) = w_0 + w_1 * x_1 + \dots + w_m * x_m$$

20


## Logistic regression - examples



- The blue model is the function  $\text{sigmoid}(1.5 \cdot x)$ , while the red line is the function  $\text{sigmoid}(x)$

21

## Logistic regression - examples



- The blue model is the function  $\text{sigmoid}(-10 + 1.5 \cdot x)$ , while the red line is the function  $\text{sigmoid}(-8 + x)$

22

## Logistic regression – more maths

- So, what would a cost function look like for Log.Reg?

$$J(W) = -\frac{1}{n} \sum_{i=1}^n (y_i \log \sigma(g(\mathbf{x}^i)) + (1-y_i) \log(1 - \sigma(g(\mathbf{x}^i))))$$

- Note the minus (-) sign at the beginning
- The expression after the sum is called the (log) likelihood of the parameters  $W$  (i.e., how likely does the model (parameterized by  $W$ ) correctly classify the target variable  $y$  based on the input vars  $\mathbf{x}$ )
- **Method: Maximum Likelihood Estimation (MLE)**. The mathematical intuition of the above can be found in Andrew Ng's video:

<https://www.youtube.com/watch?v=SHEPb1JHw5o>

23

## Summary

- In this mini lecture, we looked at the following:
  - The regression problem with the **linear regression** method (even though it's "linear", the model does NOT have to be a straight line; the **feature selection** problem)
  - The **cost** (or, loss) **function**
  - **Gradient Descent** (several variants)
  - Overfitting, Underfitting, Regularization
  - The classification problem with **logistic regression**
    - Sigmoid function, Maximum Likelihood Estimation

24